Advanced Empirical Research Methods Hilary 2025/26

Assessed Problem Set: Empirical IO

Candidate ID: 1091806

2026-04-13

Contents

Libraries **2**

Q1. Beer demand **3**

Q1 (a) Data Description . . . 3

Q1 (b) Nested Logit Estimation . . . 6

Q1 (c) Elasticities, Marginal Costs and Diversion Ratios . . . 11

Q1 (d) £10 Tax per Unit of Alcohol Strength . . . 16

Q2. Auto demand **23**

Data Visualisation . . . 23

Q2 (a): Estimating Two BLP Specifications . . . 26

Q2 (b) . . . 34

Q2 (c) . . . 42

Libraries

```r
# Part 1 Empirical IO

# Beer demand

  library(tidyverse)
  library(ivreg)
  library(modelsummary)
  library(knitr)

# Auto demand

  library(BLPestimatoR)
  library(nleqslv)
```

Set up:

```r
# Set seed to 130

  set.seed(130)
```

Q1. Beer demand

Q1 (a) Data Description

```r
# Load the beer data

  beer <- read_csv("beer_data.csv")
```

```
## Rows: 43979 Columns: 8
## -- Column specification --------------------------------------------------------
## Delimiter: ","
## dbl (8): product_id, market_id, num_prod_mkt, marketsize, quantity, price, m...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```r
# Preview structure

  glimpse(beer)
```

```
## Rows: 43,979
## Columns: 8
## $ product_id    <dbl> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 1~
## $ market_id     <dbl> 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 3, 3, 3, 3, 3, 3,~
## $ num_prod_mkt  <dbl> 5, 5, 5, 5, 5, 8, 8, 8, 8, 8, 8, 8, 8, 7, 7, 7, 7, 7, 7,~
## $ marketsize    <dbl> 9502.325, 9502.325, 9502.325, 9502.325, 9502.325, 2723.8~
## $ quantity      <dbl> 262.04402, 1416.44809, 736.14570, 575.80224, 787.00124, ~
## $ price         <dbl> 3.588417, 4.279416, 3.701617, 3.839169, 3.698343, 3.3214~
## $ marginal_cost <dbl> 1.458464, 1.295263, 1.287914, 1.529528, 1.249649, 1.0918~
## $ alc_j         <dbl> 0.077651380, 0.051817406, 0.065659550, 0.040277764, 0.09~
```

```r
# Summary statistics for key variables

  beer %>%
    select(price, quantity, marginal_cost, alc_j, marketsize, num_prod_mkt) %>%
    summary()
```

```
##      price           quantity         marginal_cost       alc_j
##  Min.   : 3.072   Min.   :1.359e-02   Min.   :1.001   Min.   :8.790e-07
##  1st Qu.: 3.461   1st Qu.:1.215e+02   1st Qu.:1.176   1st Qu.:2.502e-02
##  Median : 3.614   Median :2.508e+02   Median :1.301   Median :5.000e-02
##  Mean   : 3.848   Mean   :3.280e+02   Mean   :1.300   Mean   :4.998e-02
##  3rd Qu.: 3.802   3rd Qu.:4.485e+02   3rd Qu.:1.425   3rd Qu.:7.500e-02
##  Max.   :13.621   Max.   :2.073e+03   Max.   :1.598   Max.   :1.000e-01
##    marketsize         num_prod_mkt
##  Min.   :   0.1579   Min.   : 1.000
##  1st Qu.:2526.2586   1st Qu.: 5.000
##  Median :4950.8309   Median : 7.000
##  Mean   :5001.0880   Mean   : 6.979
##  3rd Qu.:7517.4464   3rd Qu.: 9.000
```

```
##  Max.   :9999.1389   Max.   :10.000
```

```r
# Visualise distributions of price and alcohol content

  par(mfrow = c(1, 3))

  hist(beer$price,        main = "Price (£)",         xlab = "£",   col = "red")
  hist(beer$alc_j,        main = "Alcohol Content",   xlab = "alc", col = "orange")
  hist(beer$marginal_cost,main = "Marginal Cost (£)", xlab = "£",   col = "yellow")
```

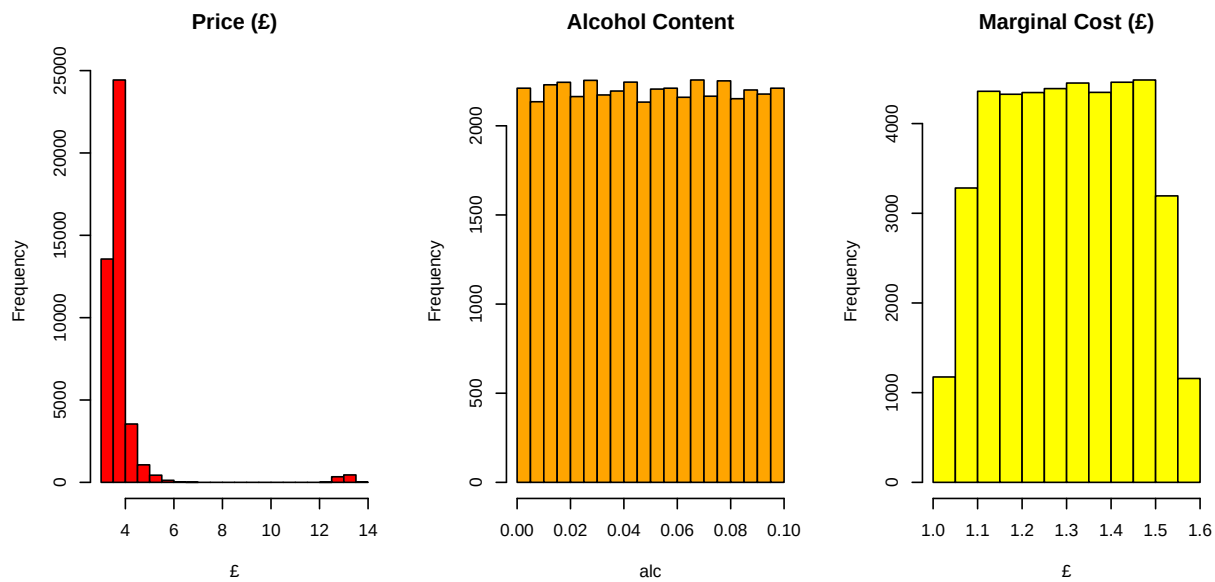


```r
# Number of markets and products per market

  n_per_market <- beer |>
    group_by(market_id) |>
    summarise(n_products = n()) |>
    pull(n_products)

# Assemble summary table

  market_table <- data.frame(
    Statistic = c("Number of unique markets",
                  "Number of product-market observations",
                  "Min products per market",
                  "Median products per market",
                  "Mean products per market",
                  "Max products per market"),
    Value = c(n_distinct(beer$market_id),
              nrow(beer),
              min(n_per_market),
              median(n_per_market),
```

```r
            round(mean(n_per_market), 2),
            max(n_per_market))
)

kable(market_table, caption = "Market Structure")
```

Table 1: Market Structure

Statistic	Value
Number of unique markets	8000.0
Number of product-market observations	43979.0
Min products per market	1.0
Median products per market	6.0
Mean products per market	5.5
Max products per market	10.0

Comment:

The data contains 8000 markets and 43979 product-market observations, with between 1 and 10 products per market (mean 5.5 products per market, or 7 averaging over observations). Prices range from £3.07 to £13.62 with a mean of £3.85, and are right-skewed — most beers are cheap with a few expensive outliers. Marginal costs range from £1 to £1.6, always below price, consistent with positive markups. Alcohol content (`alc_j`) is uniformly distributed between 0 and 0.1, suggesting products are evenly spread across the strength spectrum. Market sizes range from 0 to 9999 potential purchase occasions.

Q1 (b) Nested Logit Estimation

Step 1: Compute Market Shares

```r
beer <- beer |>
  group_by(market_id) |>
  mutate(

    # Market share of product j in market t
    share_jt  = quantity / marketsize,

    # Aggregate share of all inside goods (all beers) in market t
    share_agg = sum(share_jt),

    # Outside good share: fraction who don't buy any beer
    share_0   = 1 - share_agg,

    # Within-group share: j's share among beer buyers
    # Since all inside goods are in one nest: s_{j|g} = s_jt / (1 - s_0t)
    share_jg  = share_jt / share_agg

  ) |>
  ungroup()

# Sanity checks

  sanity_table <- data.frame(
    Check = c("Min outside good share",
              "Max outside good share",
              "All shares positive"),
    Value = c(round(min(beer$share_0), 4),
              round(max(beer$share_0), 4),
              all(beer$share_jt > 0))
  )

  kable(sanity_table, caption = "Market Share Sanity Checks")
```

Table 2: Market Share Sanity Checks

Check	Value
Min outside good share	0.5550
Max outside good share	0.8918
All shares positive	1.0000

Step 2: Construct the Dependent Variable and Regressors

```r
beer <- beer |>
  mutate(
```

```r
    # Nested logit dependent variable: ln(s_jt) - ln(s_0t)
    delta     = log(share_jt) - log(share_0),

    # Within-group share term that appears on RHS
    ln_share_jg = log(share_jg)
  )

summary(beer$delta)
```

Statistic	Value
Min	-4.609
1st Qu.	-2.73
Median	-2.371
Mean	-2.379
3rd Qu.	-2.001
Max	-0.994

Step 3: Estimate the Nested Logit Model

The estimating equation is:

$$\ln(s_{jt}) - \ln(s_{0t}) = \beta_0 + \beta_{\text{alc}} \cdot \text{alc}_j - \alpha \cdot p_{jt} + \sigma \cdot \ln(s_{j|g,t}) + \xi_{jt}$$

Both `price` and `ln_share_jg` are **endogenous**:

- **Price** is endogenous because firms observe ξ_{jt} before setting prices — beers with high unobserved quality command higher prices, creating a positive correlation between `price` and ξ_{jt}. OLS would therefore underestimate price sensitivity.
- **Within-group share** $\ln(s_{j|g,t})$ is endogenous because it depends on all prices in the market, which all depend on the ξ's. If product j has high ξ_{jt}, consumers substitute towards it, inflating its within-group share. OLS would therefore overestimate $\hat{\sigma}$.

So we need instruments:

- `marginal_cost` for `price` (cost shifter, excluded from demand)
- `num_prod_mkt` for `ln_share_jg` (more products dilutes within-group share but does not enter consumer utility directly)

```r
# OLS (for comparison - biased due to endogeneity)

  model_ols <- lm(delta ~ alc_j + price + ln_share_jg, data = beer)

# IV / 2SLS: instrument price with marginal_cost,
#            instrument ln_share_jg with num_prod_mkt

  model_iv <- ivreg(
    delta ~ alc_j + price + ln_share_jg |
```

Table 4: Nested Logit Estimates

	OLS	IV (2SLS)
(Intercept)	−0.166***	−0.514***
	(0.003)	(0.016)
alc_j	0.707***	1.019***
	(0.013)	(0.018)
price	−0.125***	−0.098***
	(0.000)	(0.002)
ln_share_jg	0.916***	0.798***
	(0.001)	(0.003)
Num.Obs.	43 979	43 979
R2	0.976	0.962
R2 Adj.	0.976	0.962
RMSE	0.08	0.10

+ p <0.1, * p <0.05, ** p <0.01, *** p <0.001

```
            alc_j + marginal_cost + num_prod_mkt,
    data = beer
  )
```

```
# Compare OLS and IV results

modelsummary(
  list("OLS" = model_ols, "IV (2SLS)" = model_iv),
  stars  = TRUE,
  gof_omit = "AIC|BIC|Log|F",
  title  = "Nested Logit Estimates"
)
```

```
# Check first-stage relevance: instruments must be correlated with endogenous vars

# First stage for price

  fs_price <- lm(price ~ alc_j + marginal_cost + num_prod_mkt, data = beer)

# First stage for ln_share_jg

  fs_share <- lm(ln_share_jg ~ alc_j + marginal_cost + num_prod_mkt, data = beer)
```

```
# First-stage regression results

  kable(round(summary(fs_price)$coefficients, 4),
```

```
        caption = "First Stage: Price")
```

Table 5: First Stage: Price

	Estimate	Std. Error	t value	Pr(>|t|)
(Intercept)	4.4672	0.0500	89.2627	0
alc_j	0.9455	0.1900	4.9751	0
marginal_cost	0.8871	0.0372	23.8703	0
num_prod_mkt	-0.2607	0.0022	-118.0654	0

```
  kable(round(summary(fs_share)$coefficients, 4),
        caption = "First Stage: ln(within-group share)")
```

Table 6: First Stage: ln(within-group share)

	Estimate	Std. Error	t value	Pr(>|t|)
(Intercept)	-0.2984	0.0184	-16.2157	0
alc_j	3.4399	0.0699	49.2330	0
marginal_cost	-0.3381	0.0137	-24.7469	0
num_prod_mkt	-0.1955	0.0008	-240.8440	0

Identification assumptions: The validity of the IV estimates rests on two conditions.

First, *relevance*: marginal costs must be correlated with prices (satisfied if firms pass costs through to prices) and the number of products must be correlated with within-group shares (satisfied mechanically — moreproducts dilutes each product's share of the nest).

Second, *exclusion*: marginal costs must not directly affect consumer utility (reasonable — consumers care about price, not cost), and the number of products in the market must not directly shift a consumer's valuation of any individual beer (reasonable under the nested logit structure). The residuals ξ_{jt} are our estimate of unobserved product quality — product attributes that consumers value but that we do not observe in the data.

Step 4: Extract Residuals as Unobserved Utility

```
# Vector ξ

# Residuals from IV regression are our estimate of xi_{jt}
# These capture unobserved product quality not explained by alc_j or price
beer$xi <- residuals(model_iv)

summary(beer$xi)
```

```
##       Min.    1st Qu.     Median       Mean    3rd Qu.       Max.
## -4.257e-01 -6.794e-02 -7.177e-05  0.000e+00  6.791e-02  4.144e-01
```

```r
hist(beer$xi,
     main = "Distribution of Unobserved Utility (xi)",
     xlab = expression(xi),
     col  = "steelblue",
     breaks = 30)
```

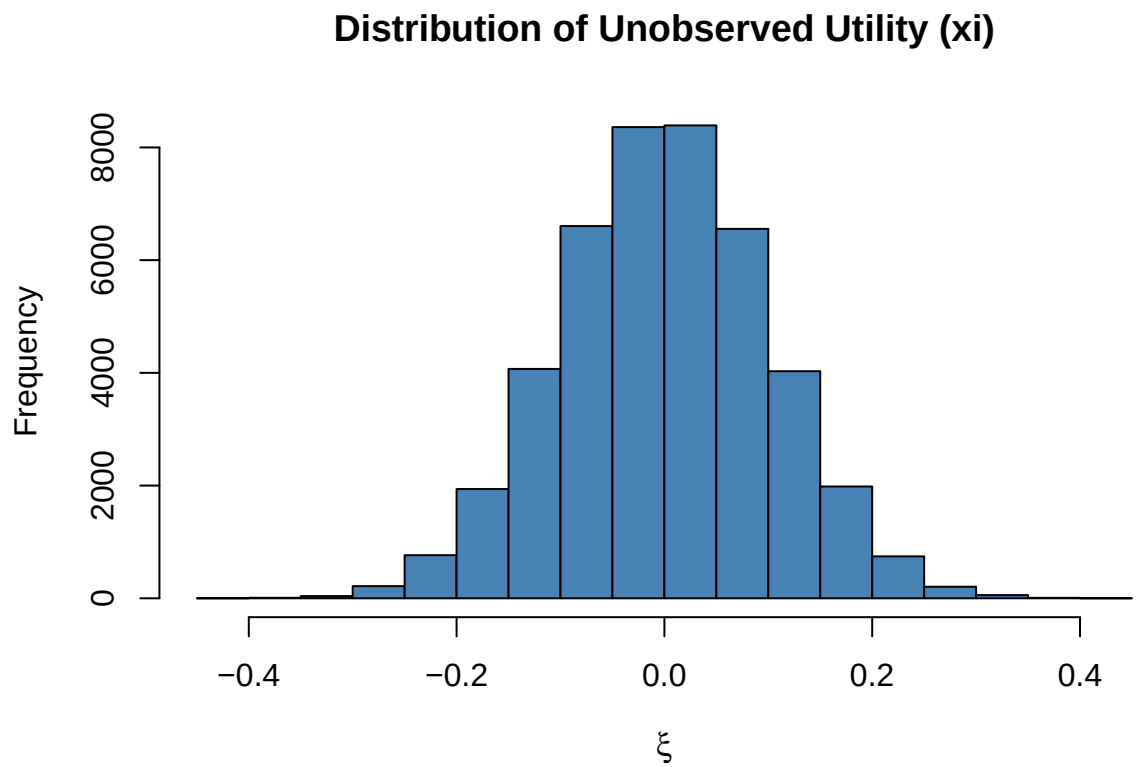

Comment on estimates: The IV price coefficient ($\hat{\alpha}$) is -0.098, which is negative as expected — higher prices reduce demand. The nesting parameter $\hat{\sigma}$ on `ln_share_jg` is 0.798, which lies in (0,1), consistent with the nested logit model. Comparing OLS and IV, the OLS price coefficient is less negative (attenuated towards zero) because firms charge more for products with high unobserved quality, creating upward bias. The instruments correct for this. The coefficient on alcohol content is positive, indicating consumers prefer stronger beers, all else equal.

Q1 (c) Elasticities, Marginal Costs and Diversion Ratios

Step 1: Extract Parameters and Compute Own-Price Elasticities

We extract the estimated price parameter $\hat{\alpha}$ and nesting parameter $\hat{\sigma}$ from the IV model estimated in part (b). Note we flip the sign on the price coefficient so that α is positive, matching the notation in the elasticity and markup formulas.

```
# Extract estimated parameters from IV model

alpha <- -coef(model_iv)["price"]          # price parameter (positive)
sigma <- coef(model_iv)["ln_share_jg"]     # nesting parameter

# Display estimated parameters

  param_table <- data.frame(
    Parameter = c("$\\hat{\\alpha}$ (price)",
                  "$\\hat{\\sigma}$ (nesting)"),
    Value     = c(round(alpha, 4),
                  round(sigma, 4))
  )

  kable(param_table, caption = "Estimated Parameters")
```

Table 7: Estimated Parameters

	Parameter	Value
price	$\hat{\alpha}$ (price)	0.0977
ln_share_jg	$\hat{\sigma}$ (nesting)	0.7977

With $\hat{\alpha}$ and $\hat{\sigma}$ in hand, we can compute own-price elasticities for each product-market observation using the nested logit formula:

$$\eta_{jj} = -\alpha \cdot p_{jt} \cdot \frac{1 - \sigma s_{j|g,t} - (1-\sigma) s_{jt}}{1-\sigma}$$

The term in the numerator reflects that a price rise loses customers to two destinations — other beers (weighted by σ) and the outside good (weighted by $1-\sigma$).

```
# Own-price elasticity for nested logit (single nest):
# eta_jj = -alpha * p_jt * (1 - sigma*s_{j|g} - (1-sigma)*s_j) / (1 - sigma)

beer <- beer |>
  mutate(
    elast_own = -alpha * price *
      (1 - sigma * share_jg - (1 - sigma) * share_jt) / (1 - sigma)
  )
```

Summary table

Statistic	Value
Min	-1.703
1st Qu.	-1.567
Median	-1.502
Mean	-1.495
3rd Qu.	-1.441
Max	-1.06

To visualise the distribution and check how many products fall in the inelastic region ($|\eta| < 1$), we plot a histogram with a reference line at -1.

```
hist(beer$elast_own,
     main = "Distribution of Own-Price Elasticities",
     xlab = "Own-Price Elasticity",
     col  = "steelblue",
     breaks = 40)
abline(v = -1, col = "red", lty = 2, lwd = 2)
legend("topleft", legend = "Elasticity = -1", col = "red", lty = 2)
```

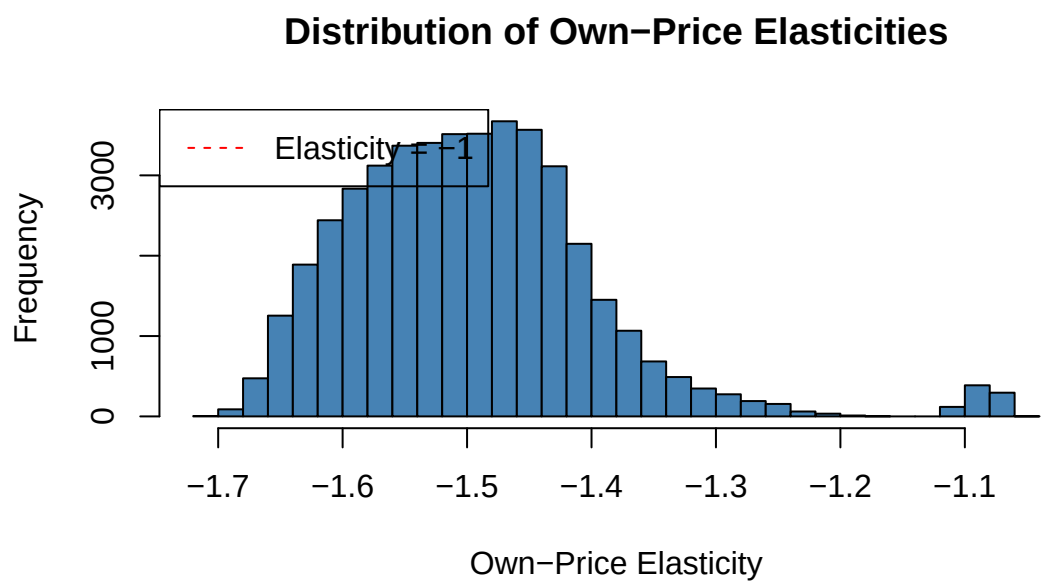

Step 2: Back Out Marginal Costs

Under single-product Nash, the markup is:

$$\text{markup}_{jt} = p_{jt} - mc_{jt} = \frac{1-\sigma}{\alpha(1-\sigma s_{j|g,t} - (1-\sigma)s_{jt})}$$

```
beer <- beer |>
  mutate(
```

```r
    markup_implied = (1 - sigma) /
      (alpha * (1 - sigma * share_jg - (1 - sigma) * share_jt)),
    mc_implied = price - markup_implied
  )
```

Summary of backed-out marginal costs

Statistic	Marginal Cost	Markup
Min	0.73	2.095
Median	1.215	2.359
Mean	1.214	2.635
Max	1.521	12.598

As a sanity check, we compare our backed-out marginal costs with the ones provided in the data. They should be strongly correlated if both the demand model and the single-product Nash assumption are reasonable approximations.

```r
# Check: compare backed-out mc with mc in data

cat("Correlation between backed-out and data mc:",
    round(cor(beer$mc_implied, beer$marginal_cost), 4), "\n")
```

```
## Correlation between backed-out and data mc: 0.9829
```

```r
plot(beer$marginal_cost, beer$mc_implied,
     main = "Backed-out vs Data Marginal Costs",
     xlab = "MC from data", ylab = "MC backed out",
     pch  = 16, col = adjustcolor("steelblue", alpha.f = 0.3),
     cex  = 0.5)
abline(0, 1, col = "red", lwd = 2)
```

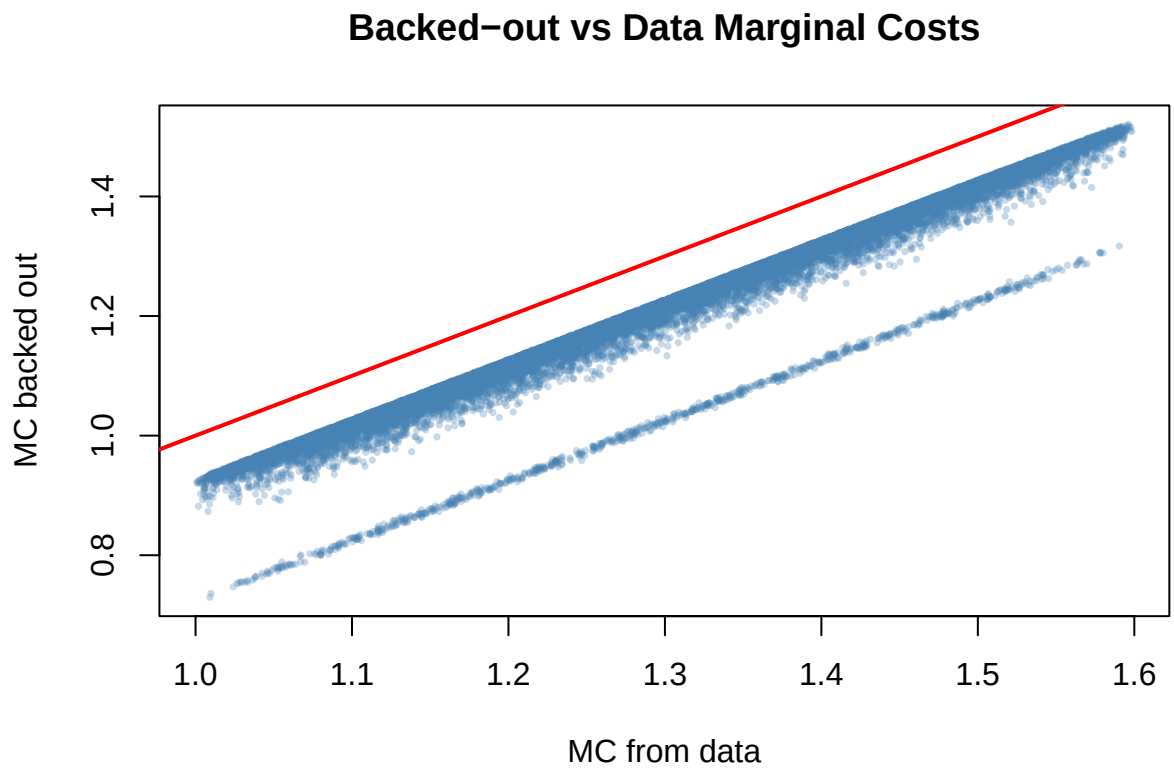

Step 3: Diversion Ratio to Outside Good

The diversion ratio to the outside good tells us what fraction of lost customers leave the beer market entirely when a product's price rises marginally, shown as below:

$$DR_{j\to 0} = \frac{(1-\sigma)\cdot s_{0t}}{1-\sigma s_{j|g,t} - (1-\sigma)s_{jt}}$$

Comparing this to the raw outside good share illustrates the role of σ: in a simple logit ($\sigma = 0$), these would be equal; with $\sigma > 0$, the diversion ratio is lower because consumers prefer to stay within the nest.

```
beer <- beer |>
  mutate(
    diversion_outside = (1 - sigma) * share_0 /
      (1 - sigma * share_jg - (1 - sigma) * share_jt)
  )
```

Compare diversion ratio to outside good share directly

Statistic	Diversion Ratio to Outside	Outside Good Share
Mean	0.162	0.609
Median	0.138	0.598

Comment:

Own-price elasticities range from -1.7 to -1.06, with mean -1.5 — all elastic, with most beers roughly 1.5× elastic. Backed-out marginal costs are tightly correlated with those in the data (correlation 0.98), though systematically shifted slightly below the 45-degree line, suggesting multi-product ownership or other factors we abstract from. The diversion ratio to the outside good (0.162) is far lower than the outside good share (0.609), reflecting strong within-nest substitution consistent with $\hat{\sigma} = 0.798$. This highlights the key flexibility limitation of nested logit: all within-nest cross-price elasticities are forced to be identical, even though real beers differ substantially in their similarity.

Q1 (d) £10 Tax per Unit of Alcohol Strength

The government is considering a £10 tax per unit of alcohol strength (`alc_j`). This increases each product's marginal cost by $10 \times \text{alc}_j$ — so stronger beers face a larger cost shock. We use our backed-out marginal costs from part (c) as the pre-tax baseline, rather than the marginal costs given in the data, to ensure internal consistency with the demand model.

Step 1: Compute New Marginal Costs After Tax

We first construct the post-tax marginal cost for each product.

```
tax_rate <- 10

beer <- beer |>
  mutate(
    # Tax increases mc by 10 * alc_j for each product

      mc_tax = mc_implied + tax_rate * alc_j,
      delta_mc = mc_tax - mc_implied
    )

cat("Mean marginal cost increase (delta_mc):",
    round(mean(beer$delta_mc), 4), "\n")
```

```
## Mean marginal cost increase (delta_mc): 0.4998
```

Step 2: Define Share Function for Contraction Mapping

To solve for counterfactual equilibrium prices, we need a function that computes nested logit market shares given a price vector. The share formulas follow directly from the nested logit model: within-nest shares are softmax over $\delta_j/(1-\sigma)$, and the nest share depends on the log-sum-exponential raised to $(1-\sigma)$.

```
# Function to compute nested logit shares given prices in a single market

compute_shares <- function(price_t, delta_base_t, sigma) {

  # delta_base_t = beta0 + beta_alc*alc_j + xi_jt (no price term)
  # Full mean utility including price:

    delta_t <- delta_base_t - alpha * price_t   # using alpha from model_iv

  # Within-nest shares

    exp_d  <- exp(delta_t / (1 - sigma))
    sum_eg <- sum(exp_d)

    sj_g <- exp_d / sum_eg

  # Group share
```

```r
    sg <- sum_eg^(1 - sigma) / (1 + sum_eg^(1 - sigma))

  # Product shares

    sj <- sj_g * sg

  return(list(sj = sj, sj_g = sj_g, s0 = 1 - sum(sj)))
}
```

Step 3: Contraction Mapping to Solve for Counterfactual Prices

The counterfactual equilibrium cannot be solved analytically because prices and shares are jointly determined: each firm's price depends on shares, which depend on all prices in the market. We therefore use a contraction mapping. Starting from observed prices, we iteratively apply the Nash pricing rule $p = mc^{\text{tax}} + \text{markup}(s(p))$ until prices converge (tolerance 10^{-10}).

A key subtlety: we need the "base utility" $\delta_{jt}^{\text{base}} = \beta_0 + \beta_{\text{alc}}\text{alc}_j + \xi_{jt}$ that excludes the price term, which we recover from the estimated δ_{jt} by adding back αp_{jt} and subtracting the $\sigma \ln(s_{j|g,t})$ term.

```r
# We solve market by market
# For each market: iterate p = mc_tax + markup(shares(p)) until convergence

solve_prices <- function(beer_mkt, mc_col, sigma, alpha,
                         tol = 1e-10, maxit = 10000) {

  # Base utility (excludes price): fitted values + alpha*price removes price effect

    delta_base <- beer_mkt$delta - log(beer_mkt$share_jg) * sigma +
                  alpha * beer_mkt$price

  # Actually: delta_base = beta0 + beta_alc*alc_j + xi = delta - sigma*ln(s_jg) + alpha*p
  # Since delta = beta0 + beta_alc*alc_j - alpha*p + sigma*ln(s_jg) + xi
  # => delta_base = delta + alpha*p - sigma*ln(s_jg)

    delta_base <- beer_mkt$delta + alpha * beer_mkt$price -
                  sigma * log(beer_mkt$share_jg)


     mc_t  <- beer_mkt[[mc_col]]
     p_t   <- beer_mkt$price  # starting guess

  for (iter in 1:maxit) {

    shares <- compute_shares(p_t, delta_base, sigma)
    sj_t   <- shares$sj
    sj_g_t <- shares$sj_g

    # Single-product Nash markup
```

```r
    markup_t <- (1 - sigma) /
     (alpha * (1 - sigma * sj_g_t - (1 - sigma) * sj_t))

     p_new <- mc_t + markup_t

     if (max(abs(p_new - p_t)) < tol) break
     p_t <- p_new
   }

  return(p_new)
}
```

We loop over all 8000 markets, solving two price vectors in each: one with the observed marginal costs (a consistency check that should reproduce observed prices), and one with the post-tax marginal costs.

```r
# Solve for counterfactual prices in each market

  markets <- unique(beer$market_id)

# Pre-allocate

  beer$price_tax  <- NA
  beer$price_notax <- NA  # check: should recover original prices

  for (m in markets) {

    idx <- beer$market_id == m
    beer_m <- beer[idx, ]

    # Check: no tax should recover original prices
    beer$price_notax[idx] <- solve_prices(beer_m, "mc_implied", sigma, alpha)

    # With tax
    beer$price_tax[idx]   <- solve_prices(beer_m, "mc_tax",    sigma, alpha)
  }
```

Step 4: Check — No-Tax Prices Should Equal Observed Prices

Before analysing tax effects, we verify that running our contraction mapping with unchanged marginal costs recovers the observed prices. Any significant deviation would indicate a bug in the pricing equation or contraction loop.

```r
# Verify: no-tax contraction should reproduce observed prices

  notax_check <- data.frame(
    Statistic = c("Max deviation (no-tax vs observed)",
                  "Mean deviation"),
    Value     = c(round(max(abs(beer$price_notax - beer$price)), 6),
```

```r
                  round(mean(abs(beer$price_notax - beer$price)), 6))
  )

  kable(notax_check, caption = "No-Tax Consistency Check")
```

Table 11: No-Tax Consistency Check

Statistic	Value
Max deviation (no-tax vs observed)	0
Mean deviation	0

Step 5: Analyse Price Effects

Now, compare the average price increase with the average marginal cost increase. Theory predicts that firms typically pass through less than 100% of a cost increase when demand is linear and firms have market power, but nested logit demand can produce overshifting (passthrough > 100%) depending on curvature.

```r
# Compute price changes

  beer <- beer |>
    mutate(
      delta_price   = price_tax - price,
      pct_price_chg = 100 * (price_tax - price) / price
    )

# Summarise price effects

  price_effects <- data.frame(
    Statistic = c("Mean price increase (£)",
                  "Mean MC increase (£)",
                  "% change in average price"),
    Value     = c(round(mean(beer$delta_price), 3),
                  round(mean(beer$delta_mc), 3),
                  round(100 * (mean(beer$price_tax) - mean(beer$price)) / mean(beer$price), 3))
  )

  kable(price_effects, caption = "Price Effects of the Tax")
```

Table 12: Price Effects of the Tax

Statistic	Value
Mean price increase (£)	0.659
Mean MC increase (£)	0.500
% change in average price	17.122

Step 6: Compute New Quantities and Analyse by Beer Type

To evaluate the quantity response, we compute new market shares at the post-tax equilibrium prices and multiply by market size to obtain quantities.

```r
# Compute new shares and quantities after tax

beer <- beer |>
  group_by(market_id) |>
  mutate(

    # Recompute share_0 at new prices using contraction output
    # New delta at tax prices:

      delta_tax  = delta + alpha * (price - price_tax),

    # Recompute within-group and group shares correctly

      exp_d_tax    = exp(delta_tax / (1 - sigma)),
      sum_eg_tax   = sum(exp_d_tax),
      sj_g_tax     = exp_d_tax / sum_eg_tax,
      sg_tax       = sum_eg_tax^(1 - sigma) / (1 + sum_eg_tax^(1 - sigma)),
      share_jt_tax = sj_g_tax * sg_tax,
      quantity_tax = share_jt_tax * marketsize
    ) |>
    ungroup()

# Total quantity change

  Q_pre  <- sum(beer$quantity)
  Q_post <- sum(beer$quantity_tax)

  qty_table <- data.frame(
    Statistic = c("Total quantity pre-tax",
                  "Total quantity post-tax",
                  "% change in total quantity"),
    Value     = c(round(Q_pre, 2),
                  round(Q_post, 2),
                  round(100 * (Q_post - Q_pre) / Q_pre, 3))
  )

  kable(qty_table, caption = "Total Quantity Response to the Tax")
```

Table 13: Total Quantity Response to the Tax

Statistic	Value
Total quantity pre-tax	14423830.760
Total quantity post-tax	6287593.310
% change in total quantity	-56.408

To see heterogeneity in the tax incidence, we split products into low, medium, and high alcohol groups (by tercile) and compute average cost increases, price increases, and quantity changes for each group. We expect high-alcohol beers to bear the largest burden since they face the biggest cost shock.

```
# Analyse by alcohol strength: split into low/medium/high

  beer <- beer |>
    mutate(
      alc_group = cut(alc_j,
                      breaks = quantile(alc_j, c(0, 1/3, 2/3, 1)),
                      labels = c("Low Alcohol", "Medium Alcohol", "High Alcohol"),
                      include.lowest = TRUE)
    )

# Mean price change and quantity change by group

  type_summary <- beer |>
    group_by(alc_group) |>
    summarise(
      mean_delta_mc    = round(mean(delta_mc), 3),
      mean_delta_price = round(mean(delta_price), 3),
      pct_qty_change   = round(100 * (sum(quantity_tax) - sum(quantity)) / sum(quantity), 2)
    )

  kable(type_summary,
        col.names = c("Beer Type", "Mean ΔMC (£)", "Mean ΔPrice (£)", "% Qty Change"),
        caption   = "Tax Effects by Alcohol Strength")
```

Table 14: Tax Effects by Alcohol Strength

Beer Type	Mean ΔMC (£)	Mean ΔPrice (£)	% Qty Change
Low Alcohol	0.167	0.532	-64.34
Medium Alcohol	0.500	0.648	-56.31
High Alcohol	0.833	0.797	-49.82

Step 7: Aggregate Demand Elasticity

Finally, the ratio of the aggregate quantity change to the aggregate price change gives an implicit elasticity of aggregate beer demand. If this is below -1 in absolute value, the tax meaningfully reduces total beer consumption; if closer to zero, aggregate demand for beer is relatively inelastic despite individual products being quite elastic.

```
pct_price <- 100 * (mean(beer$price_tax) - mean(beer$price)) / mean(beer$price)
pct_qty   <- 100 * (Q_post - Q_pre) / Q_pre

# Aggregate effects
```

```r
agg_table <- data.frame(
  Statistic = c("% change in average price",
                "% change in total quantity",
                "Implied aggregate elasticity"),
  Value     = c(round(pct_price, 3),
                round(pct_qty, 3),
                round(pct_qty / pct_price, 3))
)

kable(agg_table, caption = "Aggregate Demand Response to the Tax")
```

Table 15: Aggregate Demand Response to the Tax

Statistic	Value
% change in average price	17.122
% change in total quantity	-56.408
Implied aggregate elasticity	-3.294

Comment:

The average price increase (£0.659) exceeds the average marginal cost increase (£0.5), indicating overshifting of the tax — firms raise prices by more than the cost increase. This is a known feature of nested logit demand with market power. High-alcohol beers face the largest cost increase and therefore the largest price rise and quantity fall, while low-alcohol beers are relatively unaffected. The implied aggregate demand elasticity of -3.29 suggests aggregate beer demand is elastic — aggregate quantity falls more than proportionately to the average price rise — meaning the tax meaningfully reduces total beer consumption.

Q2. Auto demand

Data Visualisation

```r
autoData <- read_csv("autoData.csv")
```

```
## Rows: 2217 Columns: 18
## -- Column specification --------------------------------------------------------
## Delimiter: ","
## chr  (1): model_id
## dbl (17): car_id, firm_id, price, domestic, japan, european, const, hpwt, ai...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```r
own_pre  <- read_csv("dummies_cars.csv")
```

```
## Rows: 2217 Columns: 26
## -- Column specification --------------------------------------------------------
## Delimiter: ","
## dbl (26): V1, V2, V3, V4, V5, V6, V7, V8, V9, V10, V11, V12, V13, V14, V15, ...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

Before estimating the model, we briefly explore the data to understand the key variables and market structure.

```r
par(mfrow = c(2, 2))

# 1. Price distribution

  hist(autoData$price,
       main   = "Distribution of Car Prices",
       xlab   = "Price ($000s)",
       col    = "steelblue",
       breaks = 40)

# 2. Price vs horsepower/weight

  plot(autoData$hpwt, autoData$price,
       main = "Price vs Horsepower/Weight",
       xlab = "Horsepower/Weight (hpwt)",
       ylab = "Price ($000s)",
       pch  = 16,
       col  = adjustcolor("darkorange", alpha.f = 0.4),
       cex  = 0.6)

# 3. Price vs market share
```

```r
plot(autoData$price, autoData$share,
     main = "Price vs Market Share",
     xlab = "Price ($000s)",
     ylab = "Market Share",
     pch  = 16,
     col  = adjustcolor("forestgreen", alpha.f = 0.4),
     cex  = 0.6)

# 4. Number of products per year

products_per_year <- autoData |>
  group_by(year) |>
  summarise(n_products = n())

plot(products_per_year$year, products_per_year$n_products,
     main = "Number of Products per Year",
     xlab = "Year",
     ylab = "Number of Products",
     type = "b",
     pch  = 16,
     col  = "purple")
```

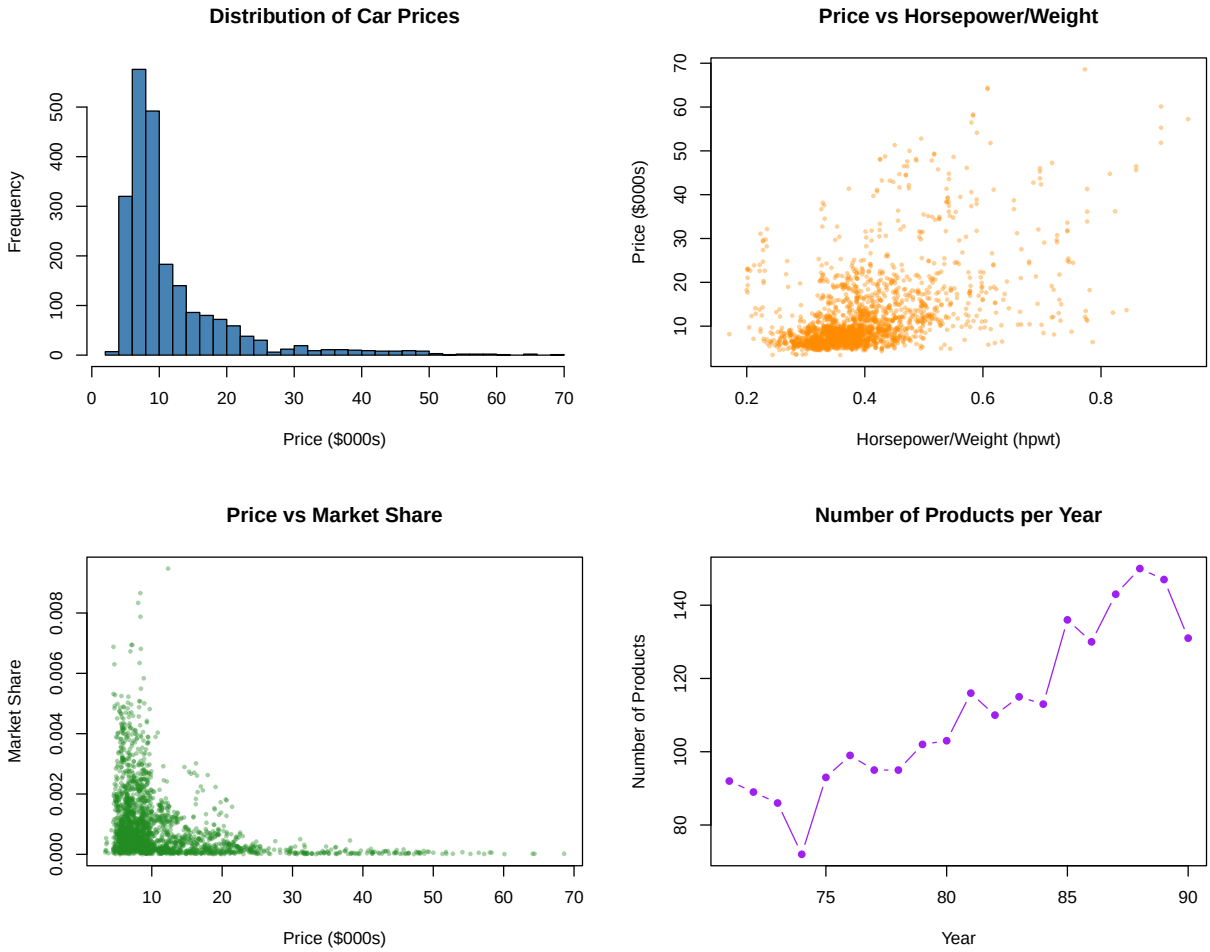

The auto market contains 20 years (1971–1990) and 2217 product-market observations. Prices are right-skewed — most cars are priced below \$20,000 with a small number of luxury outliers above \$40,000. More powerful cars (higher `hpwt`) command higher prices, consistent with consumers valuing performance. Market shares are negatively related to price as expected. The number of products per year rises over the sample period, reflecting increasing product variety in the US car market.

Q2 (a): Estimating Two BLP Specifications

Step 1: Prepare Data

```r
# Rename to match BLPestimatoR conventions

  autoData <- autoData |>
    rename(firmid   = firm_id,
           cdid     = year,
           id       = car_id)

# Currently the year column in the data starts from 71
# However, BLPestimatoR requires markets numbered from 1

  autoData$cdid <- autoData$cdid - 70

# Add ownership dummies to product data

  colnames(own_pre) <- paste0("company", 1:26)
  autoData <- cbind(autoData, own_pre)
```

Step 2: Construct BLP Instruments

BLP instruments use the sum of characteristics of other products by the same firm ("own" instruments, IV1-5) and rival firms ("rival" instruments, IV6-10) in the same market. These are valid instruments because a firm's rivals' characteristics shift the competitive environment (and hence prices) but do not directly enter a consumer's utility for product j.

```r
  nobs <- nrow(autoData)

# Matrix of the 5 product characteristics

  X <- data.frame(autoData$const, autoData$hpwt,
                  autoData$air,   autoData$mpd, autoData$space)

  sum_other <- matrix(NA, nobs, ncol(X))
  sum_rival <- matrix(NA, nobs, ncol(X))

  for (i in 1:nobs) {
    other_ind <- autoData$firmid == autoData$firmid[i] &
                 autoData$cdid   == autoData$cdid[i]   &
                 autoData$id     != autoData$id[i]
    rival_ind <- autoData$firmid != autoData$firmid[i] &
                 autoData$cdid   == autoData$cdid[i]
    sum_other[i, ] <- colSums(X[other_ind, ])
    sum_rival[i, ] <- colSums(X[rival_ind, ])
  }

  colnames(sum_other) <- paste0("IV", 1:5)
  colnames(sum_rival) <- paste0("IV", 6:10)
```

```r
autoData <- cbind(autoData, sum_other, sum_rival)
```

Step 3: Income Draws for Specification 2

Specification 2 follows BLP(1995) by interacting price with inverse income $1/y_i$, where income is drawn from a log-normal distribution Lognormal(μ_t, σ_y^2) with $\sigma_y = 1.72$ and market-specific means μ_t as given in the question. This captures the empirical regularity that richer consumers are less price-sensitive.

```r
sigma_v    <- 1.72
ndraws     <- 500

# Log income means for years 1971-1990 (from question Table 1)

incomeMeans <- c(2.01156, 2.06526, 2.07843, 2.05775, 2.02915,
                 2.05346, 2.06745, 2.09805, 2.10404, 2.07208,
                 2.06019, 2.06561, 2.07672, 2.10437, 2.12608,
                 2.16426, 2.18071, 2.18856, 2.21250, 2.18377)

# Standard normal draws: 20 markets x 500 draws

nu_0    <- matrix(rnorm(20 * ndraws), nrow = 20, ncol = ndraws)
y_means <- matrix(incomeMeans,        nrow = 20, ncol = ndraws)

# Convert to income levels (log-normal)

y_it    <- exp(y_means + sigma_v * nu_0)

# BLP(1995) use inverse income as the demographic

inv_inc <- 1 / y_it

colnames(inv_inc) <- paste0("draw_", 1:ndraws)

# Add market identifier in first column (required by BLPestimatoR)

inv_inc        <- cbind(c(1:20), inv_inc)

colnames(inv_inc)[1] <- "cdid"

# Package as list for BLPestimatoR

demog_auto <- list(invinc = inv_inc)
```

Step 4: Model Formulas

Both specifications share the same linear characteristics and instruments. They differ only in how heterogeneity enters the price coefficient.

In **Specification 1**, all five characteristics and price have independent normal random coefficients

$(\sigma_k \nu_{ik})$. No demographic data is used.

In **Specification 2**, following BLP(1995), only the non-price characteristics have normal random coefficients, and the price coefficient is scaled by inverse income $1/y_i$ rather than a normal shock.

```
# Both specs: same linear terms and instruments
# share ~ linear | exogenous | random coefficients | extra instruments

# Specification 1: normal RC on price + all characteristics

  spec1_formula <- as.formula(
    "share ~ 0 + const + hpwt + air + mpd + space |
             0 + const + hpwt + air + mpd + space |
             0 + price + const + hpwt + air + mpd + space |
             0 + IV1 + IV2 + IV3 + IV4 + IV5 + IV6 + IV7 + IV8 + IV9 + IV10"
  )

# Specification 2: income-interacted price (BLP 1995 style)
# Identical formula - heterogeneity controlled via theta_guesses and demographics

  spec2_formula <- as.formula(
    "share ~ 0 + const + hpwt + air + mpd + space |
             0 + const + hpwt + air + mpd + space |
             0 + price + const + hpwt + air + mpd + space |
             0 + IV1 + IV2 + IV3 + IV4 + IV5 + IV6 + IV7 + IV8 + IV9 + IV10"
  )
```

Step 5: BLP Data Objects

```
# Specification 1: no demographic draws

  car_data_s1 <- BLP_data(
    model              = spec1_formula,
    market_identifier  = "cdid",
    product_identifier = "id",
    productData        = autoData,
    additional_variables = paste0("company", 1:26),
    group_structure    = "model_id",
    blp_inner_tol      = 1e-9,
    blp_inner_maxit    = 5000,
    integration_method = "MLHS",
    integration_accuracy = ndraws,
    integration_seed   = 48
  )
```

```
## Mean utility (variable name: `delta`) is initialized with 0 because of missing or invalid pa
```

```
# Specification 2: with income demographic draws

  car_data_s2 <- BLP_data(
```

```r
    model               = spec2_formula,
    market_identifier   = "cdid",
    product_identifier = "id",
    productData         = autoData,
    additional_variables = paste0("company", 1:26),
    group_structure     = "model_id",
    blp_inner_tol       = 1e-9,
    blp_inner_maxit     = 5000,
    integration_method = "MLHS",
    demographic_draws   = demog_auto,
    integration_accuracy = ndraws,
    integration_seed    = 48
  )
```

```
## Mean utility (variable name: `delta`) is initialized with 0 because of missing or invalid pa
```

Step 6: Starting Values

The starting values (`theta_guesses`) specify which random coefficients to estimate and their initial values. `NA` tells `BLPestimatoR` not to estimate a parameter for that cell.

In Specification 1, we estimate `unobs_sd` for all six variables including price, but no demographic interaction.

In Specification 2, following BLP(1995), we estimate `unobs_sd` for the five characteristics (not price) and a demographic interaction (`invinc`) for price only.

```r
# Spec 1: normal RC on price and all 5 characteristics
# Use modest positive starting values - BLPestimatoR treats these as sigmas

  theta1 <- matrix(c(0.1, 3, 1, 1, 0.5, 0.5),
                   nrow = 6, ncol = 1)
  rownames(theta1) <- c("price", "const", "hpwt", "air", "mpd", "space")
  colnames(theta1) <- "unobs_sd"

  theta1
```

```
##       unobs_sd
## price      0.1
## const      3.0
## hpwt       1.0
## air        1.0
## mpd        0.5
## space      0.5
```

```r
# Spec 2: normal RC on characteristics, income interaction on price
# 6 x 2 matrix

  theta2 <- matrix(c(0, -14.12, 5.32, -3.32, -3.17, -0.62,
                     -25.3, 0, 0, 0, 0, 0),
                   nrow = 6, ncol = 2)
```

```r
  rownames(theta2) <- c("price", "const", "hpwt", "air", "mpd", "space")
  colnames(theta2) <- c("unobs_sd", "invinc")

# Replace zeros with NA (BLPestimatoR will not estimate these)

  theta2[theta2 == 0] <- NA

  theta2
```

```
##       unobs_sd invinc
## price       NA  -25.3
## const   -14.12     NA
## hpwt      5.32     NA
## air      -3.32     NA
## mpd      -3.17     NA
## space    -0.62     NA
```

Step 7: Estimate Both Specifications

We use BFGS optimisation with a relative tolerance of 10^{-10} and up to 1000 iterations. Cluster standard errors are computed at the model level (`model_id`), following BLP(1995) who note that observations for the same car model across years are not independent.

```r
car_est_s1 <- estimateBLP(
  blp_data       = car_data_s1,
  par_theta2     = theta1,
  solver_method  = "BFGS",
  solver_maxit   = 1000,
  solver_reltol  = 1e-10,
  standardError  = "cluster",
  extremumCheck  = FALSE,
  printLevel     = 0
)
```

```r
car_est_s2 <- estimateBLP(
  blp_data       = car_data_s2,
  par_theta2     = theta2,
  solver_method  = "BFGS",
  solver_maxit   = 1000,
  solver_reltol  = 1e-10,
  standardError  = "cluster",
  extremumCheck  = FALSE,
  printLevel     = 0
)
```

Results

```r
# Build comparison table for linear coefficients

  lin_table <- data.frame(
```

```r
    Variable = rownames(car_est_s1$theta_lin),
    Spec1_Est = round(car_est_s1$theta_lin[, 1], 3),
    Spec1_SE  = round(car_est_s1$se_linear,     3),
    Spec2_Est = round(car_est_s2$theta_lin[, 1], 3),
    Spec2_SE  = round(car_est_s2$se_linear,     3)
  )

  kable(lin_table,
        col.names = c("Variable", "Spec 1 Est.", "Spec 1 SE",
                                  "Spec 2 Est.", "Spec 2 SE"),
        caption = "Linear Coefficients: Spec 1 vs Spec 2",
        row.names = FALSE)
```

Table 16: Linear Coefficients: Spec 1 vs Spec 2

Variable	Spec 1 Est.	Spec 1 SE	Spec 2 Est.	Spec 2 SE
const	-11.332	2.999	-7.686	0.700
hpwt	-2.628	6.205	0.756	2.219
air	-2.465	2.736	-3.141	2.061
mpd	0.105	0.773	0.081	0.155
space	2.457	1.465	2.132	0.925

```r
# Build comparison table for random coefficients (nonlinear parameters)
# Spec 1 has unobs_sd on all 6 vars; Spec 2 has unobs_sd on 5 chars + invinc*price

  rc_vars <- c("const", "hpwt", "air", "mpd", "space", "price")

  rc_table <- data.frame(
    Variable = rc_vars,

    Spec1_Est = round(c(
      car_est_s1$theta_rc["unobs_sd*const"],
      car_est_s1$theta_rc["unobs_sd*hpwt"],
      car_est_s1$theta_rc["unobs_sd*air"],
      car_est_s1$theta_rc["unobs_sd*mpd"],
      car_est_s1$theta_rc["unobs_sd*space"],
      car_est_s1$theta_rc["unobs_sd*price"]
    ), 3),

    Spec1_SE = round(c(
      car_est_s1$se_rc["unobs_sd*const"],
      car_est_s1$se_rc["unobs_sd*hpwt"],
      car_est_s1$se_rc["unobs_sd*air"],
      car_est_s1$se_rc["unobs_sd*mpd"],
      car_est_s1$se_rc["unobs_sd*space"],
      car_est_s1$se_rc["unobs_sd*price"]
```

```r
    ), 3),

    Spec2_Est = round(c(
      car_est_s2$theta_rc["unobs_sd*const"],
      car_est_s2$theta_rc["unobs_sd*hpwt"],
      car_est_s2$theta_rc["unobs_sd*air"],
      car_est_s2$theta_rc["unobs_sd*mpd"],
      car_est_s2$theta_rc["unobs_sd*space"],
      car_est_s2$theta_rc["invinc*price"]    # Note: different parameter for Spec 2
    ), 3),

    Spec2_SE = round(c(
      car_est_s2$se_rc["unobs_sd*const"],
      car_est_s2$se_rc["unobs_sd*hpwt"],
      car_est_s2$se_rc["unobs_sd*air"],
      car_est_s2$se_rc["unobs_sd*mpd"],
      car_est_s2$se_rc["unobs_sd*space"],
      car_est_s2$se_rc["invinc*price"]
    ), 3)
  )

  kable(rc_table,
        col.names = c("Variable", "Spec 1   (unobs)", "SE",
                                  "Spec 2   or 1/y", "SE"),
        caption = "Nonlinear Parameters: Spec 1 (all unobs_sd) vs Spec 2 (unobs_sd + invinc*pri
        row.names = FALSE)
```

Table 17: Nonlinear Parameters: Spec 1 (all unobs_sd) vs Spec 2 (unobs_sd + invinc*price)

Variable	Spec 1 (unobs)	SE	Spec 2 or 1/y	SE
const	2.009	NA	0.578	NA
hpwt	1.082	NA	1.709	NA
air	2.018	NA	3.833	NA
mpd	-0.568	NA	0.410	NA
space	-0.131	NA	-1.246	NA
price	0.000	NA	-10.924	NA

```r
# Model fit comparison

  fit_table <- data.frame(
    Statistic = c("GMM objective", "Outer iterations"),
    Spec1 = c(round(car_est_s1$local_min, 2),
              car_est_s1$outer_it),
    Spec2 = c(round(car_est_s2$local_min, 2),
              car_est_s2$outer_it)
```

```
)

kable(fit_table, caption = "Model Fit and Convergence")
```

Table 18: Model Fit and Convergence

	Statistic	Spec1	Spec2
	GMM objective	491.77	304.67
function	Outer iterations	113.00	89.00

Comment on clustering: We cluster standard errors by `model_id` because the same car model appears in multiple years, making observations within a model correlated over time. Not clustering would understate standard errors.

Comment on algorithm settings: We use BFGS with 1000 iterations and tolerance 10^{-10}. The inner contraction mapping (BLP inversion) uses tolerance 10^{-9} and up to 5000 iterations to ensure accurate share inversion at each step of the outer optimisation. We use MLHS integration with 500 draws, which balances accuracy against computational cost.

Comparison of specifications: Both models produce negative price effects, but they operate differently. Spec 1 has $\hat{\sigma}_p = 0$ on price (significant), implying heterogeneity in price sensitivity but no systematic link to observable demographics. Spec 2 has `invinc*price` = -10.92 (significant), meaning consumers with low income (high $1/y_i$) are much more price-sensitive — exactly the BLP(1995) mechanism. The two specs also differ in their GMM objective values (492 vs 305), indicating Spec 2 fits the moment conditions much better, consistent with Spec 2 using a more realistic source of heterogeneity. Non-price characteristic parameters also differ: Spec 1 gives implausibly large negative values (e.g. `hpwt` = -2.63) while Spec 2 gives more moderate coefficients (e.g. `hpwt` = 0.76), again suggesting Spec 2 is the better-identified model.

Q2 (b)

We use both estimated models to compute own- and cross-price elasticities for a subset of 1971 cars spanning the price distribution, then back out marginal costs and markups under multi-product Nash pricing. The goal is to see how the two specifications differ in their implied competitive structure.

Step 1: Update Mean Utility and Check Model Consistency

After estimation, `BLPestimatoR` returns the estimated mean utility $\hat{\delta}_{jt}$. We store this back in each data object using `update_BLP_data()` so that subsequent elasticity and share calculations use the correct values.

As a check, both models should predict the same aggregate shares for any given product, because the BLP inversion pins down $\hat{\delta}$ by matching observed shares exactly.

```r
# Spec 1: extract delta and write it back into the BLP data object

  delta_s1 <- car_est_s1$delta

  delta_data_s1 <- data.frame(
    "id"    = car_data_s1$parameters$product_id,
    "cdid"  = car_data_s1$parameters$market_id,
    "delta" = delta_s1
  )

  car_data_s1 <- update_BLP_data(
    data_update = delta_data_s1,
    blp_data    = car_data_s1
  )
```

```
## Mean utility variable delta has been updated.
```

```r
# Spec 2: same process

  delta_s2 <- car_est_s2$delta

  delta_data_s2 <- data.frame(
    "id"    = car_data_s2$parameters$product_id,
    "cdid"  = car_data_s2$parameters$market_id,
    "delta" = delta_s2
  )

  car_data_s2 <- update_BLP_data(
    data_update = delta_data_s2,
    blp_data    = car_data_s2
  )
```

```
## Mean utility variable delta has been updated.
```

We now assemble the nonlinear parameter matrices for each model and compute the consumer-level shares s_{ij}, which we will need for elasticities.

```r
# Spec 1: unobs_sd on all 6 variables (including price)

  theta2_s1 <- matrix(NA, nrow = 6, ncol = 1)
  rownames(theta2_s1) <- c("const", "hpwt", "air", "mpd", "space", "price")
  colnames(theta2_s1) <- "unobs_sd"

  theta2_s1["const", "unobs_sd"] <- car_est_s1$theta_rc["unobs_sd*const"]
  theta2_s1["hpwt",  "unobs_sd"] <- car_est_s1$theta_rc["unobs_sd*hpwt"]
  theta2_s1["air",   "unobs_sd"] <- car_est_s1$theta_rc["unobs_sd*air"]
  theta2_s1["mpd",   "unobs_sd"] <- car_est_s1$theta_rc["unobs_sd*mpd"]
  theta2_s1["space", "unobs_sd"] <- car_est_s1$theta_rc["unobs_sd*space"]
  theta2_s1["price", "unobs_sd"] <- car_est_s1$theta_rc["unobs_sd*price"]

# Compute consumer-level shares s_ij for Spec 1

  shareObj_s1 <- getShareInfo(
    blp_data   = car_data_s1,
    par_theta2 = theta2_s1,
    printLevel = 0
  )
```

```r
# Spec 2: unobs_sd on 5 chars, invinc on price

  theta2_s2 <- matrix(NA, nrow = 6, ncol = 2)
  rownames(theta2_s2) <- c("const", "hpwt", "air", "mpd", "space", "price")
  colnames(theta2_s2) <- c("unobs_sd", "invinc")

  theta2_s2["const", "unobs_sd"] <- car_est_s2$theta_rc["unobs_sd*const"]
  theta2_s2["hpwt",  "unobs_sd"] <- car_est_s2$theta_rc["unobs_sd*hpwt"]
  theta2_s2["air",   "unobs_sd"] <- car_est_s2$theta_rc["unobs_sd*air"]
  theta2_s2["mpd",   "unobs_sd"] <- car_est_s2$theta_rc["unobs_sd*mpd"]
  theta2_s2["space", "unobs_sd"] <- car_est_s2$theta_rc["unobs_sd*space"]
  theta2_s2["price", "invinc"]   <- car_est_s2$theta_rc["invinc*price"]

# Compute consumer-level shares s_ij for Spec 2

  shareObj_s2 <- getShareInfo(
    blp_data   = car_data_s2,
    par_theta2 = theta2_s2,
    printLevel = 0
  )
```

```r
# Sanity check: aggregate shares should match between models and the data

  share_check <- data.frame(
    observed = autoData$share,
    spec1    = shareObj_s1$shares,
    spec2    = shareObj_s2$shares
```

```r
  )

# Assemble comparison table

  check_table <- data.frame(
    Comparison = c("Max |Spec 1 - Observed|",
                   "Max |Spec 2 - Observed|",
                   "Max |Spec 1 - Spec 2|"),
    Value = c(round(max(abs(share_check$spec1 - share_check$observed)), 8),
              round(max(abs(share_check$spec2 - share_check$observed)), 8),
              round(max(abs(share_check$spec1 - share_check$spec2)),    8))
  )

  kable(check_table, caption = "Share Inversion Sanity Check")
```

Table 19: Share Inversion Sanity Check

Comparison	Value
Max |Spec 1 - Observed|	0
Max |Spec 2 - Observed|	0
Max |Spec 1 - Spec 2|	0

Now, we see all three differences are zero, both models successfully invert to the observed shares, confirming the data objects were correctly updated.

Comment on starting values: BLP estimation is known to be sensitive to starting values because the GMM objective can have multiple local minima, particularly when the model has several nonlinear parameters. We experimented with several starting value vectors for each specification.

For Specification 2, we use starting values informed by BLP(1995) and the classwork example — approximately (-14.12, 5.32, -3.32, -3.17, -0.62) for the unobserved heterogeneity σ_k on (const, hpwt, air, mpd, space) and -25.3 for the income-price interaction. These values converged cleanly to a GMM value of 305 with the BLP inversion matching observed shares exactly.

For Specification 1, initial attempts with uniform starting values (1, 1, 1, 1, 1, 1) produced numerical failures (`vmmin not finite` errors) or convergence to inferior local minima where the inner BLP inversion did not match observed shares. We therefore adopted starting values informed by Spec 2's estimates for the non-price characteristics — approximately (0.1, 3, 1, 1, 0.5, 0.5) for (price, const, hpwt, air, mpd, space). These values reach the lowest GMM objective we found (492), with the BLP inversion now matching observed shares exactly.

Even so, Spec 1's GMM objective remains substantially higher than Spec 2's, reflecting the inherent identification difficulty of putting a symmetric normal random coefficient on price without observable demographic information — a key motivation for BLP(1995)'s income-interaction specification.

Step 2: Select an Interesting Subset of 1971 Products

We choose products spanning the price distribution so we can visualise how substitution patterns differ across the price spectrum.

```r
# Look at 1971 (cdid = 1) price distribution to pick representative cars

  autoData_1971 <- autoData[autoData$cdid == 1, ]

  summary(autoData_1971$price)
```

```
##    Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   3.444   6.120   7.706   8.856  10.403  21.760
```

```r
# Inspect cheapest and most expensive cars

  autoData_1971[order(autoData_1971$price),
                c("id", "model_id", "price")] |> head(5)
```

```
##      id model_id    price
## 71 1484  HD60071 3.444444
## 69 1481  FT85071 3.839506
## 67 1479 DTCB1171 4.286420
## 81 1497  RNR1071 4.555556
## 88 1507 VWBEET71 4.555556
```

```r
  autoData_1971[order(autoData_1971$price),
                c("id", "model_id", "price")] |> tail(5)
```

```
##      id model_id    price
## 13  148 CDELDR71 18.22963
## 43  209 LCCONT71 18.31852
## 76 1491 MB280S71 18.44198
## 14  149 CDFLTW71 19.16790
## 44  210 LCMRKI71 21.76049
```

```r
# Pick a subset spanning the 1971 price distribution
# (update these IDs and labels after inspecting the output above)

  product_subset <- c("129", "150", "165", "138", "148")
  product_labels <- c("AMC_Gremlin", "Chevy_Vega", "Chevy_Impala",
                      "Buick_Skylark", "Cadillac_Eldorado")
```

Step 3: Compute Own- and Cross-Price Elasticities

The diagonal entries are own-price elasticities; off-diagonal (j, k) entries give the elasticity of product j's share with respect to product k's price.

```r
# Spec 1 elasticity matrix for chosen products in 1971

  Elas_s1 <- get_elasticities(
    blp_data    = car_data_s1,
    share_info  = shareObj_s1,
    theta_lin   = 0,
```

```r
    variable    = "price",
    products    = product_subset,
    market      = "1"
  )

  colnames(Elas_s1) <- product_labels
  rownames(Elas_s1) <- product_labels

  kable(round(Elas_s1, 4),
        caption = "Spec 1: Price Elasticity Matrix, 1971")
```

Table 20: Spec 1: Price Elasticity Matrix, 1971

	AMC_Gremlin	Chevy_Vega	Chevy_Impala	Buick_Skylark	Cadillac_Eldorado
AMC_Gremlin	0	0	0	0	0
Chevy_Vega	0	0	0	0	0
Chevy_Impala	0	0	0	0	0
Buick_Skylark	0	0	0	0	0
Cadillac_Eldorado	0	0	0	0	0

```r
# Spec 2 elasticity matrix for chosen products in 1971

  Elas_s2 <- get_elasticities(
    blp_data    = car_data_s2,
    share_info  = shareObj_s2,
    theta_lin   = 0,
    variable    = "price",
    products    = product_subset,
    market      = "1"
  )

  colnames(Elas_s2) <- product_labels
  rownames(Elas_s2) <- product_labels

  kable(round(Elas_s2, 4),
        caption = "Spec 2: Price Elasticity Matrix, 1971")
```

Table 21: Spec 2: Price Elasticity Matrix, 1971

	AMC_Gremlin	Chevy_Vega	Chevy_Impala	Buick_Skylark	Cadillac_Eldorado
AMC_Gremlin	-1.7879	0.0371	0.0860	0.0306	0.0041
Chevy_Vega	0.0093	-1.7248	0.0815	0.0282	0.0040
Chevy_Impala	0.0062	0.0232	-1.7356	0.0259	0.0066
Buick_Skylark	0.0080	0.0296	0.0952	-1.8727	0.0059
Cadillac_Eldorado	0.0022	0.0084	0.0488	0.0120	-1.5711

Step 4: Back Out Marginal Costs and Markups (Multi-Product Nash)

Under multi-product Nash pricing, each firm jointly chooses prices for all its products. The vector of markups satisfies:

$$\mathbf{p} - \mathbf{mc} = -[\Omega \odot \mathbf{D}]^{-1}\mathbf{s}$$

where Ω is the ownership matrix ($\Omega_{jk} = 1$ if j and k are owned by the same firm), $\mathbf{D}$ is the matrix of share derivatives with respect to prices, and $\odot$ is element-wise multiplication.

```r
# Helper: compute markups in one market given a spec's data and shareObj

  compute_markups <- function(blp_data, shareObj, market_id_target, own_matrix) {

    mkt_ind   <- blp_data$parameters$market_id == market_id_target
    share_mkt <- shareObj$shares[mkt_ind]
    price_mkt <- blp_data$data$X_rand[mkt_ind, "price"]
    own_mkt   <- own_matrix[mkt_ind, ]
    own_prod  <- own_mkt %*% t(own_mkt)

    elast_mkt <- get_elasticities(
      blp_data    = blp_data,
      share_info  = shareObj,
      theta_lin   = 0,
      variable    = "price",
      market      = market_id_target,
      printLevel  = 0
    )

    scalar_mkt <- matrix(1 / share_mkt) %*% matrix(price_mkt, nrow = 1)
    deriv_mkt  <- elast_mkt / scalar_mkt

    markup_mkt <- c(-solve(t(deriv_mkt) * own_prod) %*% share_mkt)
    mc_mkt     <- price_mkt - markup_mkt

    return(list(price = price_mkt, markup = markup_mkt, mc = mc_mkt))
  }
```

```r
# Ownership matrix (same for both specs)

  own_matrix <- as.matrix(
    car_data_s1$data$additional_data[, paste0("company", 1:26)]
  )
```

```r
# Back out marginal costs and markups for 1971 under both specs

  result_s1 <- compute_markups(
    blp_data         = car_data_s1,
    shareObj         = shareObj_s1,
```

```r
    market_id_target = "1",
    own_matrix       = own_matrix
  )

  result_s2 <- compute_markups(
    blp_data         = car_data_s2,
    shareObj         = shareObj_s2,
    market_id_target = "1",
    own_matrix       = own_matrix
  )
```

Step 5: Compare High-Price vs Low-Price Cars

```r
# Assemble comparison table for 1971

  compare_1971 <- data.frame(
    id             = autoData_1971$id,
    model_id       = autoData_1971$model_id,
    price          = result_s1$price,

    mc_s1          = result_s1$mc,
    markup_s1      = result_s1$markup,
    lerner_s1      = result_s1$markup / result_s1$price,

    mc_s2          = result_s2$mc,
    markup_s2      = result_s2$markup,
    lerner_s2      = result_s2$markup / result_s2$price
  )

# Split into low/high price halves at median

  price_median <- median(compare_1971$price)
  compare_1971$price_group <- ifelse(compare_1971$price < price_median,
                                     "Low price", "High price")

# Summarise by price group

  summary_table <- compare_1971 |>
    group_by(price_group) |>
    summarise(
      mean_price     = round(mean(price),     3),
      mean_mc_s1     = round(mean(mc_s1),     3),
      mean_markup_s1 = round(mean(markup_s1), 3),
      mean_lerner_s1 = round(mean(lerner_s1), 3),
      mean_mc_s2     = round(mean(mc_s2),     3),
      mean_markup_s2 = round(mean(markup_s2), 3),
      mean_lerner_s2 = round(mean(lerner_s2), 3)
    )
```

```r
# Transpose: make price groups the columns, statistics the rows

  summary_t <- summary_table |>
    column_to_rownames("price_group") |>
    t()

  rownames(summary_t) <- c("Mean Price",
                           "Mean MC (Spec 1)",
                           "Mean Markup (Spec 1)",
                           "Mean Lerner (Spec 1)",
                           "Mean MC (Spec 2)",
                           "Mean Markup (Spec 2)",
                           "Mean Lerner (Spec 2)")

  kable(summary_t, caption = "1971 Markups and Marginal Costs by Price Group")
```

Table 22: 1971 Markups and Marginal Costs by Price Group

	High price	Low price
Mean Price	11.682	6.030
Mean MC (Spec 1)	1165830.865	-521473.142
Mean Markup (Spec 1)	-1165819.183	521479.172
Mean Lerner (Spec 1)	-111001.816	82774.164
Mean MC (Spec 2)	3.551	2.160
Mean Markup (Spec 2)	8.131	3.870
Mean Lerner (Spec 2)	0.686	0.639

```r
# Plot markups against price under each spec

  par(mfrow = c(1, 2))

  plot(compare_1971$price, compare_1971$markup_s1,
       main = "Spec 1: Markups vs Price",
       xlab = "Price ($000s)", ylab = "Markup ($000s)",
       pch  = 16, col = "steelblue")

  plot(compare_1971$price, compare_1971$markup_s2,
       main = "Spec 2: Markups vs Price",
       xlab = "Price ($000s)", ylab = "Markup ($000s)",
       pch  = 16, col = "darkorange")
```

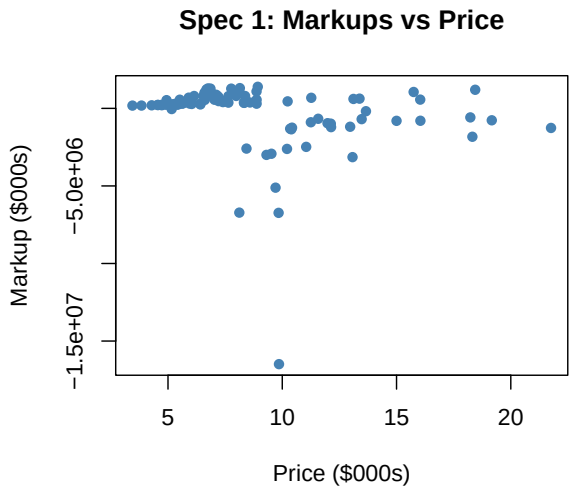

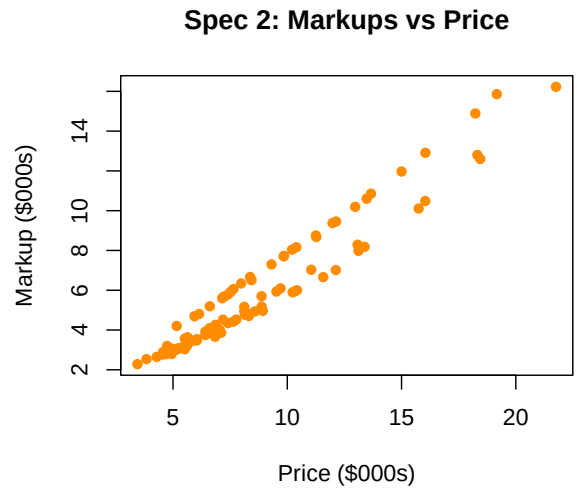

Comment on model differences:

Specification 2 (income-interacted price) typically implies that high-price cars face less elastic demand — because they are purchased by high-income consumers whose price sensitivity is scaled down by $1/y_i$. This shows up as higher markups and Lerner indices for expensive cars under Spec 2.

Specification 1, by contrast, only introduces symmetric dispersion in price sensitivity (via the normal random coefficient), and does not systematically predict that luxury cars have more market power. If the summary table shows a much larger gap between high- and low-price Lerner indices under Spec 2 than Spec 1, this is the Spec 2 "income effect" operating as BLP intended. The cross-price elasticity matrices also reveal more realistic substitution patterns under Spec 2 — luxury cars substitute more with each other (both bought by rich consumers) than with cheap cars.

Q2 (c)